

Tried on Ubuntu 16.04

### **STEP 1 : Installation**

- Open up terminal
- ``mkdir tools && cd tools``
- ``wget http://apache.claz.org/hbase/stable/hbase-1.2.6-bin.tar.gz``
- ``tar -xvzf hbase-1.2.6-bin.tar.gz``
- ``cd hbase-1.2.6/``

For HBase 0.98.5 and later, we are required to set the JAVA\_HOME environment variable before starting Hbase using conf/hbase-env.sh.

The JAVA\_HOME variable should be set to a directory which contains the executable file bin/java.

- ``sudo gedit ~/.profile``
- Add these two lines at the end

```
export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-amd64
export PATH=$PATH:$JAVA_HOME/bin
```
- ``source ~/.profile``
- ``echo $JAVA_HOME`` to check if the path is set correctly

Now starting the shell

The ./bin/start-hbase.sh script is provided as a convenient way to start HBase.

- ``./bin/start-hbase.sh``  
starting master, logging to  
/home/rshah9/tools/hbase-1.2.6/bin/./logs/hbase-rshah9-master-bn16-246.dcs.mcnc.org.out  
OpenJDK 64-Bit Server VM warning: ignoring option PermSize=128m; support was removed in 8.0  
OpenJDK 64-Bit Server VM warning: ignoring option MaxPermSize=128m; support was removed in 8.0
- ``./bin/hbase shell``  
thbase(main):001:0>

Now the shell has started.

- Inside the shell type ``version`` to check version or type ``help``. You can exit using ``exit`` command.

### **STEP 2: Basic Queries**

- Connect to your running instance of HBase using the hbase shell command.
- Use the create command to create a new table. You must specify the table name and the Column-Family name.

```
hbase> create 'test', 'cf'
0 row(s) in 1.2200 seconds
```

- Use the list command to see the List Information About your Table.

```
hbase> list 'test'
TABLE
test
1 row(s) in 0.0350 seconds
=> ["test"]
```

- To put data into your table, use the put command.

```
hbase> put 'test', 'row1', 'cf:a', 'value1'
0 row(s) in 0.1770 seconds
```

- Use the scan command to scan the table for data.

```
hbase> scan 'test'
ROW          COLUMN+CELL
row1         column=cf:a, timestamp=1403759475114,    value=value1
1 row(s) in 0.0440 seconds
```

- To get a single row of data at a time, use the get command.

```
hbase> get 'test','row1'
COLUMN      CELL
cf:a        timestamp=1403759475114, value=value1
1 row(s) in 0.0230 seconds
```

- If you want to delete a table or change its settings, you need to disable the table first, using the disable command. You can re-enable it using the enable command.

```
hbase> disable 'test'
0 row(s) in 1.6270 seconds
```

- hbase> enable 'test'
- ```
0 row(s) in 0.4500 seconds
```

- To drop (delete) a table, use the drop command.

```
hbase> drop 'test'
0 row(s) in 0.2900 seconds
```

- To exit the HBase Shell type ``exit``

- use `./bin/stop-hbase.sh` script to stop hbase service  
`$ ./bin/stop-hbase.sh`  
 stopping hbase..... \$

### STEP 3: HBase Thrift

- ``cd ..`` (Now you are in tools directory)
- ``sudo apt install python-dev``
- ``sudo apt update && sudo apt upgrade && sudo apt autoremove`` (to update all packages, will take some time)
- ``sudo apt-get install libboost-dev libboost-test-dev libboost-program-options-dev libboost-filesystem-dev libboost-thread-dev libevent-dev automake libtool flex bison pkg-config g++ libssl-dev``
- ``wget https://archive.apache.org/dist/thrift/0.6.0/thrift-0.6.0.tar.gz``
- ``tar xzf thrift-0.6.0.tar.gz``
- ``cd thrift-0.6.0/``
- ``./configure``
- ``sudo make`` (Will take some time)
- ``sudo make install``
- ``thrift`` to see if it is working
- ``cd ../``
- Download Hbase.thrift file from here <https://drive.google.com/file/d/0B5dejdhAYHztOGt6OGs5ZTZ3WEk/view?usp=sharing> and put it in current folder (``pwd`` for current directory)
- ``thrift -gen py Hbase.thrift`` (This will create a folder named gen-py)
- Start the hbase service if you stopped it in step 1 (inside tools/hbase-1.2.6/ do ``./bin/start-hbase.sh``)
- ``./hbase-1.2.6/bin/hbase thrift start``
- Open new terminal
- ``cd tools/gen-py``
- ``gedit table.py`` and paste below code (excluding the first line and quotes at end)  
 table.py (See explanation below)  

```

python
from thrift.transport.TSocket import TSocket
from thrift.transport.TTransport import TBufferedTransport
from thrift.protocol import TBinaryProtocol
from hbase import Hbase

transport = TBufferedTransport(TSocket('127.0.0.1', 9090))
transport.open()
protocol = TBinaryProtocol.TBinaryProtocol(transport)
client = Hbase.Client(protocol)
print(client.getTableNames())

```

...

- ``sudo apt install python-pip``
- ``sudo pip install thrift``
- ``python table.py`` should give output of the tables that you have created in the hbase database.  
['test']

Explanation of the code that we just wrote :

The below commands will import all the required HBase Thrift modules.

```
from thrift.transport.TSocket import TSocket
from thrift.transport.TTransport import TBufferedTransport
from thrift.protocol import TBinaryProtocol
from hbase import Hbase
```

The below command will create the socket transport and line protocol and allows the Thrift client to connect and talk to the Thrift server.

```
transport = TBufferedTransport(TSocket('localhost', 9090))
```

Next we need to open the socket to the Thrift server.

```
transport.open()
```

Tbinary is binary implementation of thrift (converting transport to binary implementation)

```
protocol = TBinaryProtocol.TBinaryProtocol(transport)
```

The below lines create the Client object which will be used to interact with HBase. From this client object, you will issue all your Gets and Puts.

```
client = Hbase.Client(protocol)
print(client.getTableNames())
```

#### ***STEP 4: HBase Complex Application with APIs***

We will be creating a todo application using Hbase and Python.

- ``cd ~/tools``
- ``git clone https://github.com/riken Shah/flask-hbase-todos.git``
- ``cd flask-hbase-todos``
- ``sudo pip install happybase flask``

- In new terminal open hbase shell as described in STEP 1
- In the shell do ``create 'todos', 'todos'``
- In another shell start thrift server like in Step 3. (``./tools/hbase-1.2.6/bin/hbase thrift start``)
- In current terminal do ``python app.py``
- In browser visit ``http://127.0.0.1:8080/``.
- Explore the code and play around
- Inside of third terminal, do check how entries are made from shell (Use ``scan`` command as described in Step 2)

The main logic resides in ``todoapp.py`` file. It has following code.

```
``python
from datetime import datetime
import happybase
import json
import time
import calendar
from flask import Flask, request, flash, url_for, redirect, \
    render_template, abort

app = Flask(__name__)
app.config.from_pyfile('todoapp.cfg')
time_format = "%Y-%m-%d %H:%M:%S"
time_format_w_ms = "%Y-%m-%d %H:%M:%S.%f"

connection = happybase.Connection(app.config['HBASE_HOST'], app.config['HBASE_PORT'])
column_and_key = app.config['HBASE_TABLE']+":"+app.config['HBASE_TABLE']

@app.route('/')
def index():
    return render_template('index.html', todos=getTasks())

@app.route('/new', methods=['GET', 'POST'])
def new():
    if request.method == 'POST':
        if not request.form['title']:
            flash('you must provide a basic description', 'error')
        elif not request.form['text']:
            flash('additional notes are required', 'error')
        else:
            saveTask({"title": request.form['title'], "text":request.form['text'], "pub_date": str(datetime.now()),
"done": "False" })
            flash(u'Todo item was successfully created')
            return redirect(url_for('index'))
```

```

return render_template('new.html')

@app.route('/todos/<int:id>', methods = ['GET' , 'POST'])
def show_or_update(id):
    item = getTask(id)
    if request.method == 'GET':
        return render_template('view.html',todo=item)
    item['title'] = request.form['title']
    item['pub_date'] = item['pub_date'].strftime(time_format_w_ms)
    item['text'] = request.form['text']
    item['done'] = str(('done.%d' % id) in request.form)
    if request.method == 'POST':
        saveTask(item)
    return redirect(url_for('index'))

def getTasks():
    hbase = connection.table(app.config['HBASE_TABLE'])
    print "get all"
    results = []
    for key, data in hbase.scan():
        print key, data
        results.append(_decode(data))
    return results

def getTask(id):
    print "get "+str(id)
    #timestamp = _unix_to_datetime(id)
    hbase = connection.table(app.config['HBASE_TABLE'])
    #return _decode(hbase.row(timestamp))
    return _decode(hbase.row(str(id)))

def saveTask(task):
    item = _encode(task)
    pub_date = datetime.strptime(task['pub_date'], time_format)
    id = str(_datetime_to_unix(pub_date))
    hbase = connection.table(app.config['HBASE_TABLE'])
    hbase.put(id, {column_and_key: str(item)})

def delTask(id):
    hbase = connection.table(app.config['HBASE_TABLE'])
    table.delete(str(id))
    print "deleting: "+str(id)

```

```

def resetTable():
    dropTable()
    createTable()

def dropTable():
    print "dropping our table..."
    if( app.config['HBASE_TABLE'] in connection.tables()):
        connection.delete_table( app.config['HBASE_TABLE'], True)

def createTable():
    a = {app.config['HBASE_TABLE']: dict()}
    print a
    if( app.config['HBASE_TABLE'] not in connection.tables()):
        print "creating a table schema..."
        connection.create_table( app.config['HBASE_TABLE'], a)

def _encode(item):
    pub_date = datetime.strptime(item['pub_date'], time_format_w_ms)
    item['pub_date'] = pub_date.strftime(time_format)
    item['done'] = str(item['done'])
    return json.dumps(item)

def _decode(item):
    task = json.loads(item[column_and_key])
    pub_date = datetime.strptime(task['pub_date'], time_format)
    done = True if task['done'] == 'True' else False
    id = _datetime_to_unix(pub_date)
    return {"title": str(task['title']), "text": str(task['text']), "pub_date": pub_date, "done": done, "id": id }

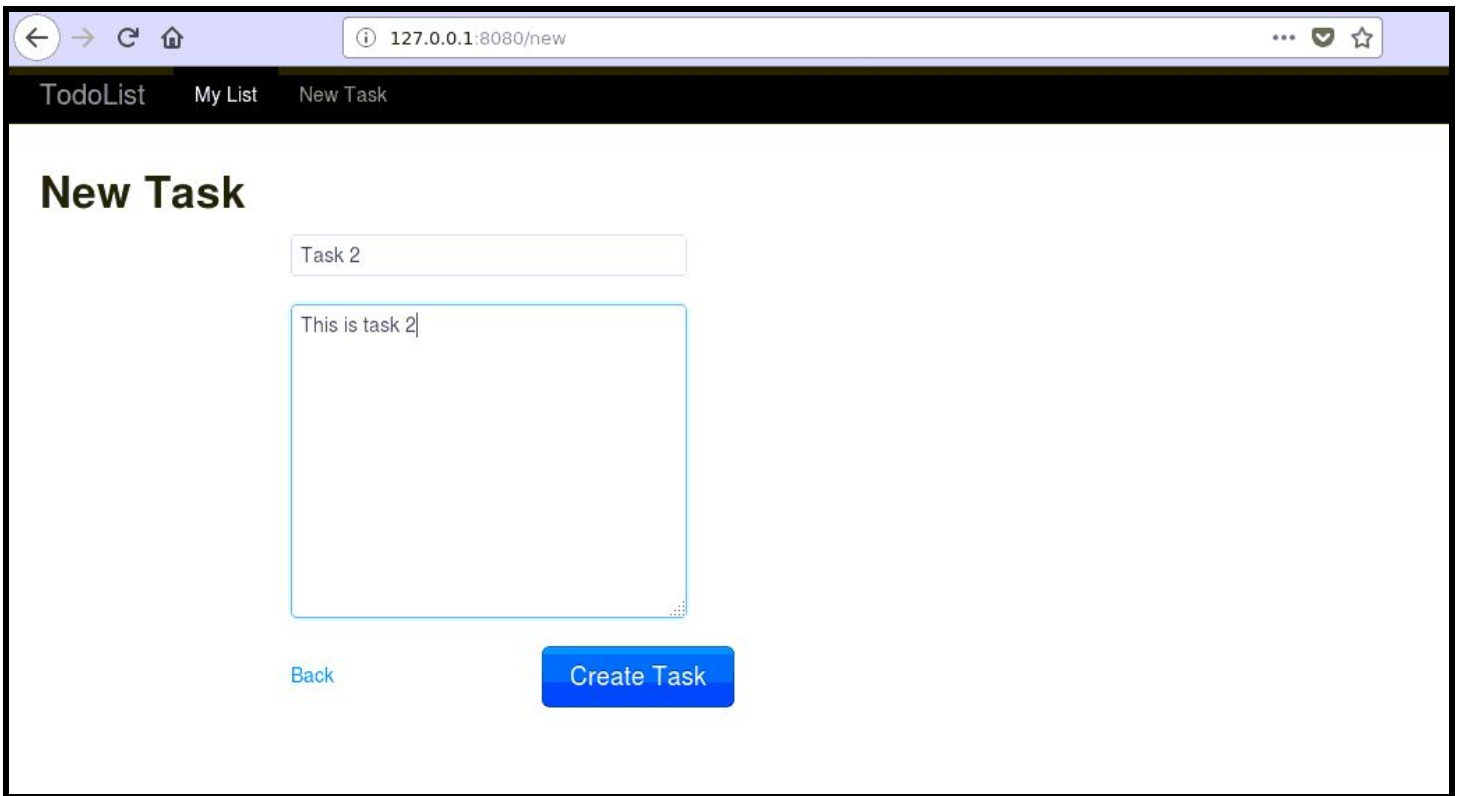
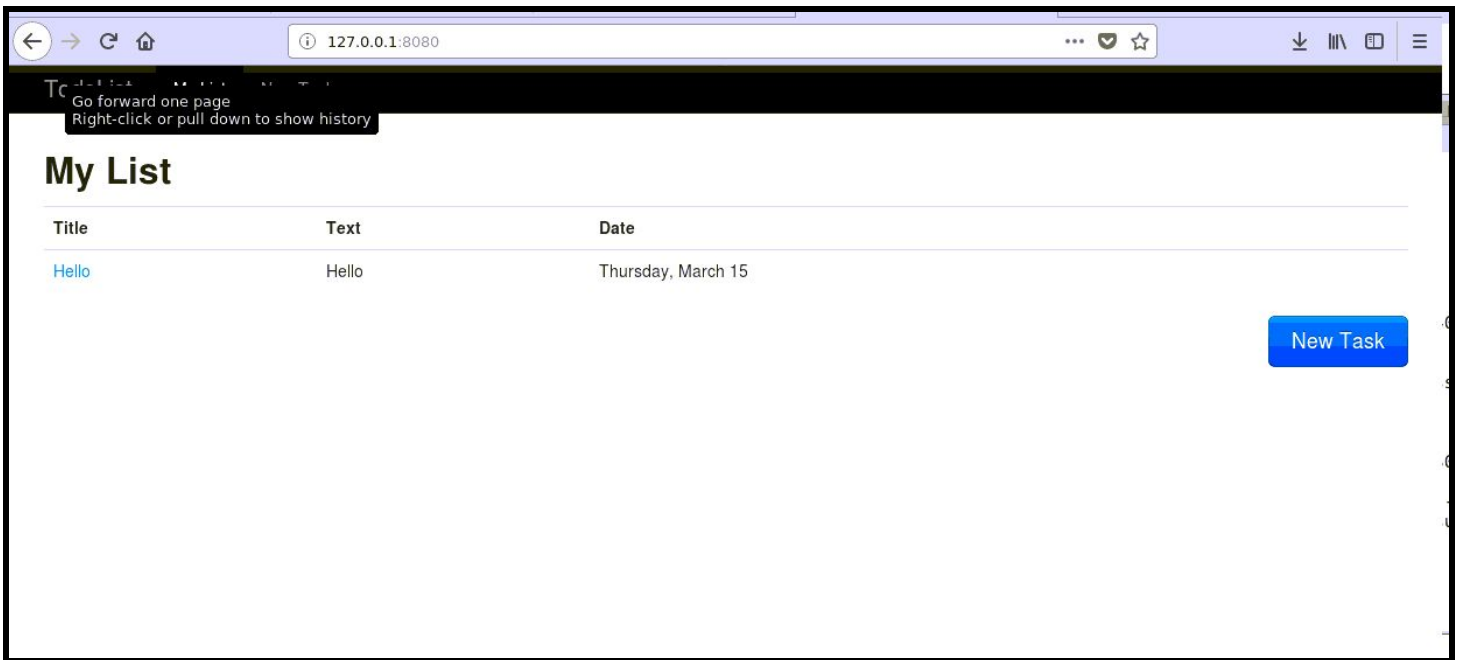
def _unix_to_datetime(id):
    return datetime.fromtimestamp(int(id)+25200).strftime(time_format)

def _datetime_to_unix(date):
    return calendar.timegm(date.timetuple())

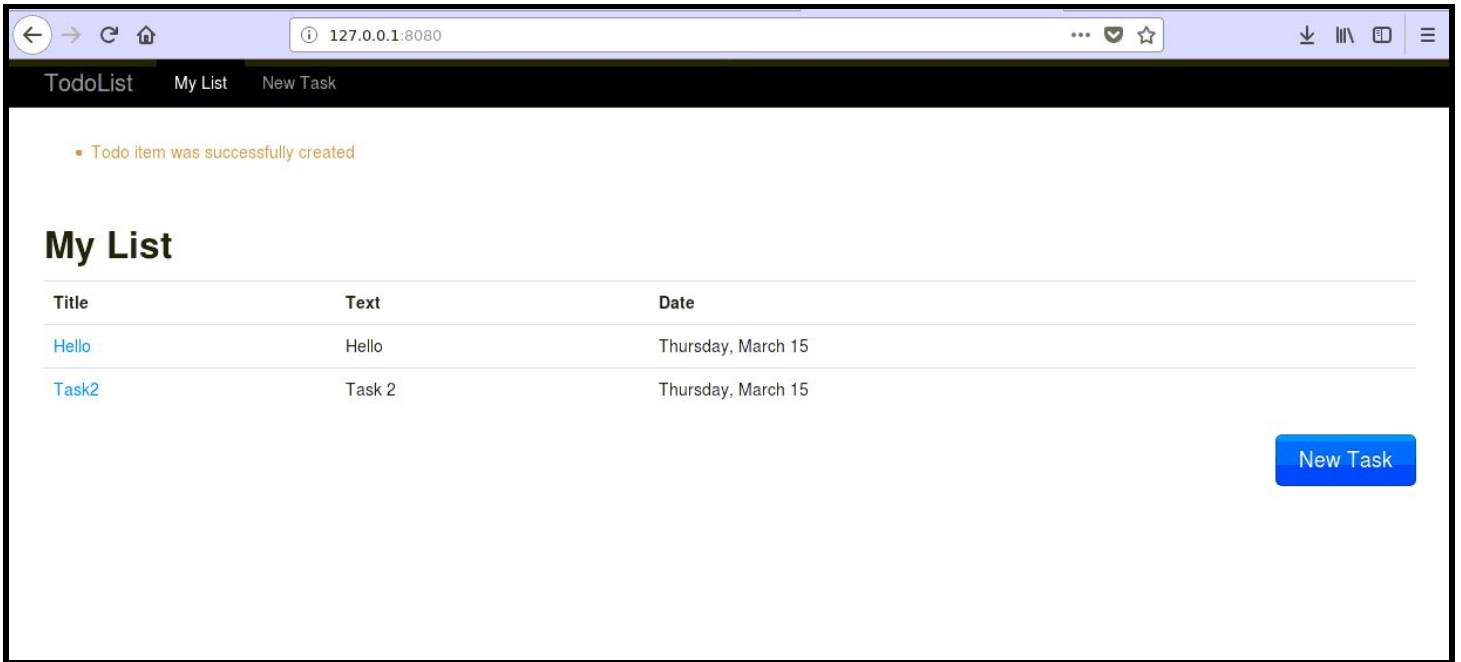
if __name__ == '__main__':
    print connection.tables()
    #createTable()
    resetTable()
    app.run()

```

...







THANK YOU...!